

whitespace

AI LEAD GENERATION FOR DUPONT TEDLAR · GRAPHICS & SIGNAGE

Six agents take public trade show data and produce a ranked list of sales prospects, each with a named decision maker and an outreach draft that has been checked against its own evidence before it can be sent. One run costs about \$11 and reruns are free.

RESULTS FROM ONE RUN

Companies sourced from exhibitor directories	819
Removed by keyword screen, at zero cost	83
Triaged by model knowledge, no web search	736 for \$0.10
Enriched with sourced firmographics	116 of 120
Qualified after rubric scoring	92
Disqualified with a stated reason	28
Qualified leads with a named decision maker	12 of 92
Outreach drafts written	26
Drafts passing the evidence check	21 of 26
Total API cost	about \$11
Cost per qualified lead, researched and drafted	\$0.42
Equivalent manual effort	about 2 working days
Cost of a full rerun	\$0, served from cache

THE ONE RESULT THAT MAKES THE CASE

3M Commercial Graphics scores 94 on fit and 12 on leverage. A textbook ideal customer, and an account that cannot be won, because 3M manufactures its own protective overlaminates. A single qualification score puts them at the top of the call list. Two scores put them in the displacement quadrant, and the outreach agent writes them a different message: it leads on a technical comparison and closes with *no need to jump on a call*, rather than asking for a meeting it will not get.

Researching one company by hand, checking fit, finding the decision maker and writing a personalized note runs 30 to 45 minutes. This does not replace that person: it does the research and the first draft and hands over a reviewable queue with the evidence attached, so the time goes into judgement and conversation instead of tab-switching. Nothing sends itself.

HOW THE WORK IS SPLIT

Stage	In	Out	What it costs	Why
Source	—	819	free	Exhibitor directories, read through one adapter per platform
Screen	819	736	free	Regex removes printers, inks, software, associations, with a reason
Classify	736	736	\$0.10	Model recall, no search: places ORAFOL and Drytac that regex cannot
Enrich	120	116	\$4.02	Web search for firmographics; every field carries a source URL
Score	116	92	\$3.06	Model reads evidence, Python applies weights from config
Contacts	92	12	\$0.60	Named people from public pages only, never inferred
Outreach	92	26	\$2.20	Draft, then a separate pass verifies every claim

Four decisions that shaped the system

1 · Spend the cheapest resource that can answer the question

Running web search on all 819 companies would have cost more than the entire project. Instead each stage is cheaper than the one it feeds. Regex removes what a pattern can settle. The model then places names regex cannot read: *ORAFOL Americas* and *Drytac* say nothing about films, and both are core prospects. Only the survivors reach the stage that pays for search. Ordering the pipeline by cost is what makes it affordable at eight hundred companies and still affordable at eight thousand.

2 · The model gathers evidence, Python computes the score

For each rubric dimension the model picks a level off a scale and quotes the fact behind it. It never sees a weight and never produces a total. The weights live in `config/icp.yaml` and the arithmetic happens in code. Three things follow: reruns are stable, every score decomposes into named dimensions a rep can audit, and retuning the rubric is a text edit rather than eight hundred model calls. Asking a model for a score out of 100 returns a number that feels right and cannot be defended.

3 · Nothing enters the dataset without a source

Every value carries a URL, a timestamp, and a confidence. A figure the model produced without a URL is dropped rather than downgraded. Where sources disagree, precedence decides: a company's own site outranks a provider record, which outranks an aggregator. That rule earned itself immediately, since four listing sites publish four different date ranges for the same Dubai trade show. Provenance is also what makes the outreach check possible at all.

4 · Two axes, because one number hides the question that matters

Fit asks whether Tedlar should want a company. Leverage asks whether they can win it, from four public signals: whether the company names a protective film partner, whether its published outdoor-life claims fall short of what a premium film allows, whether it is visibly sourcing materials now, and how established a global supplier already is in its home market. The quadrant a lead lands in decides the sales motion, and the outreach agent writes accordingly.

Leverage is deliberately held at lower confidence than fit. Finding no named film partner is weak evidence, because plenty of companies never publish supplier relationships, and the dashboard shows that rather than hiding it.

VERIFICATION, THE PART THAT DECIDES WHETHER THIS IS USABLE

Fluent personalized email is easy to generate and easy to get wrong. A note confidently referencing a product line a company does not make is worse than a generic one: it is the moment a prospect concludes nobody did the work. So drafting is two calls. The writer lists every factual claim it made. A second call, with no memory of writing the draft, checks each claim against the evidence bundle. A writer checking its own work agrees with itself, which is the whole reason for the separation.

Unsupported claims block the send and appear in the dashboard. In this run the gate caught a draft to Mactac asserting a move toward sustainable substrates that nothing in the evidence supported. If the verifier fails to respond, the draft is flagged rather than passed: silence is not approval.

Scaling, integrations, and what this does not do

SCALING

Adding an event is a block in `config/events.yaml`. Adding a *platform* is one file in `src/sources/`. ISA Sign Expo and PRINTING United are different shows run by different associations in different cities, and they cost the same as one show because both run on MapYourShow, along with several hundred other North American trade shows. Retargeting to a different customer is one config file: the product description, the disqualifiers, the rubric, the weights, and the titles worth finding all live there.

Reading MapYourShow took a decision worth naming. Its exhibitor pages are client-rendered and contain no company names at all. The working route is the print export, which is server-rendered and complete. It was chosen over the internal JSON endpoint because a print view exists for people who print things and therefore has a reason to keep working, while undocumented endpoints get renamed.

INTEGRATIONS

Every data source sits behind one interface, so swapping the free web provider for a paid one changes nothing else. **Clay** is specified with its field map and the four steps to make it live; it would beat the web provider on firmographics for private companies, which is the weakest part of the free path, and its precedence is already resolved below a company's own published figures.

LinkedIn Sales Navigator has no open API. Access runs through the partner programme, a reseller, or Clay's LinkedIn integration. So the pipeline never depends on it: names come from public pages, and the Sales Navigator link is *constructed* from the company and the target titles in config. A rep with a seat clicks into the right filtered search; a rep without one still has a name and a title. Designing around the real constraint beat stubbing an endpoint that does not exist.

ASSUMPTIONS

The brief does not define the ICP. It gives one example company, Avery Dennison Graphics Solutions, and five reasons it qualifies. The ICP here is derived from that example plus what Tedlar physically is: a protective overlamine sold into the product lines of companies that manufacture graphic films, print media, and outdoor fabrics. Not sign shops, and not distributors. Five of the six fit dimensions are those five reasons. The sixth, environmental severity, is added: Tedlar only wins where cheaper laminates fail, so conditions in the markets a company serves predict willingness to pay. It scores markets served, not headquarters.

As a check, the brief's own reference customer was run through the finished rubric. Avery Dennison scores in the top tier, which does not prove the rubric is right but does show it agrees with the one worked example available.

LIMITS

Named contacts were found for 12 of the 26 qualified leads, and the system never invents the rest. An empty contact is a correct answer, and a plausible fabricated one is the most damaging output this could produce, because a rep would act on it. Association member directories were evaluated as a second source type: ISA's is inside a members-only portal that companies must opt into, so exhibitor lists remain the only public route into this industry.

India was selected as a second region on data availability and then dropped as a source. The directories exist and are public, but render client-side and return nothing, and the exhibitor mix skews to equipment vendors the ICP filters out. Availability and relevance are different tests and only the first had been run. The events remain in config so the run reports the coverage gap rather than the region quietly disappearing, and the severity dimension survived unchanged because it always scored markets served rather than company location.

One run of 120 enrichments was lost to an unguarded merge step that raised on a headcount published as a range. Two structural changes came out of it: results are journaled to disk as they arrive, and every model response is cached, so the pipeline now replays end to end for nothing.